Epistemic Role Framing Elicits Deep Collaboration in LLM Multi-Agent Systems

Anonymous Author(s)
Affiliation
Address
email

Abstract

Evaluating whether LLM agents genuinely collaborate—rather than merely exchange outputs—requires metrics beyond task accuracy. We introduce the Collaborative Phase Profile (CPP), a PISA-grounded framework measuring collaborative depth, quality, Phase-A exploration, and collaborative yield, and the CIDI benchmark: 140 MATH problems with information splits designed to induce epistemic interdependence. Across 1,967 conversations on CIDI-amenable problems, active-learner framing is associated with a large CDI gain ($0.362 \rightarrow 0.613$, $d = 1.30$, $p < 0.001$) and more than doubles fully coupled collaboration (11% to 24%), while peer-awareness alone yields no practically meaningful gain after correction. Two full-corpus ablations show that neither active-learner framing without the split (C_{FULL}) nor expert framing with the split (C_{EXP}) reproduces this pattern, consistent with a synergy between structural interdependence and communicative role framing. Agent Epistemic Contribution (AEC) separates structural interdependence from activated epistemic necessity: 85% of problems are both-solo-zero, yet only 25% become activated under the strongest tested framing (upper-bound estimate; see §4.1). We release CPP, CIDI splits, conversation logs, scoring scripts, AEC scripts, and a benchmark card documenting corpus conditionality and intended use.

1 Introduction

Multi-agent systems built on LLMs have demonstrated remarkable capabilities in code generation [Hong et al., 2023], debate-based reasoning [Liang et al., 2023, Du et al., 2023], and socially grounded interaction [Park et al., 2023]. In most of these systems, collaboration is designed at the *task level*: agents receive different tools, different roles in a pipeline, or access to different data shards. Relatively little attention has been paid to whether the *epistemic stance* an agent adopts—how it positions itself cognitively with respect to its own knowledge and its partner's knowledge—can alter the quality of emergent collaboration.

This question has both evaluation and design implications. On the evaluation side, benchmarking collaborative agents requires metrics that can distinguish genuine joint reasoning from efficient answer-delegation—a distinction that task-accuracy scores obscure. On the design side, research in the learning sciences has long established that productive collaboration requires participants to engage in genuine co-construction of understanding rather than answer-sharing or expertise-transfer [Roschelle and Teasley, 1995, Dillenbourg, 1999]. If LLM agents default to the role of a knowledgeable assistant—which is precisely what instruction fine-tuning optimises for—eliciting collaborative behaviour requires an explicit epistemic intervention, not merely a task decomposition. Developing and validating such interventions demands evaluation infrastructure that does not yet exist for LLM multi-agent systems.

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

We operationalise this intervention through three increasingly specific framing conditions (see Section 4). The weakest framing (C2) provides no special instructions beyond handing each agent a partial information packet. The intermediate framing (C6) additionally makes agents *aware* that a peer with complementary information exists. The strongest framing (C7) instructs agents to adopt the communicative stance of an active student: express uncertainty explicitly, ask clarifying questions, verify understanding aloud before advancing to solution steps, and treat the partner's contributions as material to reason about.

To quantify how deeply agents collaborate, we introduce the Collaborative Phase Profile (CPP) framework (Section 3), grounded in the PISA 2015 CPS competency model [OECD, 2015]. CPP decomposes each conversation into a 4×3 matrix of collaborative phases and competencies and derives four summary statistics: CDI (breadth of phase engagement), CQI (quality-weighted depth), PHAQ (quality of joint exploration), and CY (composite yield).

Contributions.

1. **CPP**: a PISA-grounded 12-cell rubric producing CDI, CQI, PHAQ, CY, and a quadrant taxonomy (COUPLING / PROD_FAIL / TRIVIAL / COLLAPSE).
2. **CIDI benchmark**: 140 MATH problems with epistemic-interdependence splits verified by AEC (85% both-solo-zero).
3. **Framing study** (1,967 + 832 ablation conversations): active-learner framing is associated with a large CDI gain ($d = 1.30$); peer-awareness alone is not; neither component alone suffices.
4. **AEC decomposition**: separates structural (85% both-solo-zero) from activated (25% epistemic necessity) interdependence.

2 Background

The PISA 2015 CPS framework characterises collaborative problem solving along a 4×3 grid of process phases (Exploring, Representing, Planning, Monitoring) and collaboration competencies (Shared Understanding, Action, Team Organisation) [OECD, 2015, Graesser et al., 2017, 2018], with OECD-based tools validated in human-to-agent settings [Rojas et al., 2021]. We adapt this grid to LLM-to-LLM dialogues as the basis for CPP (Section 3).

Prior multi-agent LLM work has explored role-playing cooperation [Li et al., 2023, Hong et al., 2023, Qian et al., 2023], multi-agent conversation [Wu et al., 2024], debate-based reasoning [Liang et al., 2023, Du et al., 2023], and social-intelligence evaluation [Park et al., 2023, Zhou et al., 2023]. Debate systems optimise adversarial refinement; social-agent benchmarks assess behaviour; neither operationalises CPS phases or epistemic interdependence—the distinction CPP captures as *coordination* vs. *collaboration*.

Epistemic interdependence—where no participant can reach the goal alone—predicts productive collaboration in human learning [Barron, 2003, Johnson and Johnson, 2009], but capable LLMs may bypass it by inferring missing information. Communicative role framing (learner vs. expert stance) activates different discourse strategies in peer tutoring [Chi et al., 2001, King, 1992]; whether it does so in LLM-to-LLM settings is untested. CIDI and C7 operationalise and test both mechanisms jointly.

3 The CPP Framework

3.1 The 12-Cell Collaborative Phase Matrix

We adapt the PISA CPS 4×3 grid to LLM-to-LLM mathematical dialogues. Each cell (P, C) for process $P \in \{A, B, C, D\}$ and competency $C \in \{1, 2, 3\}$ is annotated by a trained scorer on a 0–3 scale using a detailed rubric (see Appendix A): **0** = absent, **1** = attempted but superficial, **2** = substantive but incomplete, **3** = full and well-articulated. A cell is considered *active* if its score ≥ 1.

3.2 Summary Metrics

Collaborative Depth Index (CDI).

$$\text{CDI} = \frac{|\{(P, C) : \text{score}(P, C) \geq 1\}|}{12} \in [0, 1]. \tag{1}$$

CDI measures *breadth*: what fraction of the 12 cells were genuinely engaged. A value of 1.0 means every phase-competency combination was visited; 0.0 means the dialogue was entirely unilateral.

Collaborative Quality Index (CQI).

$$\text{CQI} = \frac{\sum_{P,C} \text{score}(P, C)}{36} \in [0, 1]. \tag{2}$$

CQI weights by quality: 36 is the maximum possible total score (12×3). Unlike CDI, CQI penalises superficial engagement.

Phase A Quality (PHAQ).

$$\text{PHAQ} = \frac{\sum_{C \in \{1,2,3\}} \text{score}(A, C)}{9} \in [0, 1]. \tag{3}$$

Process A is the locus of joint problem-model construction; PHAQ $= 0$ indicates agents proceeded to solution without any joint exploration.

ATC$_{\text{CQI}}$ (ATC_CQI). We also score against the five ATC21S collaborative competency dimensions [Griffin et al., 2012] (0–3 each); aggregate $\text{ATC_CQI} = \frac{\sum 5 \text{ dims}}{15}$.

Collaborative Yield (CY).

$$\text{CY} = 0.35 \cdot \text{CDI} + 0.45 \cdot \mathbb{1}[\text{correct}] + 0.20 \cdot \delta_{\text{couple}}, \tag{4}$$

where $\delta_{\text{couple}} = 1$ iff the conversation is in the COUPLING quadrant.

3.3 Quadrant Taxonomy

We embed conversations in a 2D outcome space defined by CDI (process) and binary correctness (outcome):

Quadrant	CDI	Correct
COUPLING	≥ 0.5	Yes
PROD_FAIL	≥ 0.5	No
TRIVIAL	< 0.5	Yes
COLLAPSE	< 0.5	No

COUPLING (high CDI, correct) is the ideal—deep collaboration *and* success. PROD_FAIL (high CDI, incorrect) is productive failure [Kapur, 2016]. TRIVIAL (low CDI, correct) implies one agent solved independently. COLLAPSE (low CDI, incorrect) is complete failure of both process and outcome.

3.4 The CIDI Information Split

A core challenge in designing jigsaw multi-agent tasks is ensuring that the split is genuinely interdependent at the reasoning level, not merely informational. We introduce Collaborative Interdependence-Directed Information (CIDI): an automated pipeline that uses GPT-4.1 to decompose each problem into three components:

1. *Shared context*: the problem statement elements visible to both agents, including notation and domain.

2. *Packet A*: information given exclusively to Agent A, chosen so that Agent A cannot solve the problem without the *reasoning output* (not just the raw facts) of Agent B.
3. *Packet B*: the symmetric complement for Agent B.

The key criterion is *reasoning-output dependency*: neither agent can progress past initial exploration without the other's deductions. Unlike negative-instruction ("constitutional") splits—which LLMs frequently override using general domain knowledge—CIDI uses positive, non-inferable complementary constraints.

4 Experimental Design

4.1 Problem Corpus Selection (Phase 1)

We drew problems from the MATH benchmark [Hendrycks et al., 2021], spanning six subject areas (Algebra, Geometry, Number Theory, Prealgebra, Precalculus, and Counting & Probability) and five difficulty levels (MATH levels 1–5), yielding a $6 \times 5 = 30$ cell grid. From an initial pool of 300 problems we applied a *genuine epistemic filter*: we ran all 300 problems under condition C7 and retained only those for which $\text{CDI} \geq 0.5$ in at least two of three replications. The purpose of this filter is to identify problems for which the CIDI split produces genuine reasoning-level interdependence—problems that are too easy or too structurally simple tend to be resolved in one agent's opening message regardless of framing, making them uninformative for studying collaborative dynamics. The filter retained 140 problems (47% retention rate), distributed across all 30 subject-difficulty cells with a minimum of 1 problem per cell and a mean of 4.7 problems.

Selection conditionality. Because the filter uses C7 performance to define the corpus, *all estimates are conditional on* CIDI-*amenable problems*: absolute CDI values reflect performance on problems pre-selected for collaborative potential, and within-corpus effect sizes may overestimate the benefit expected on an unfiltered sample. Paired comparisons remain internally valid within this regime. As a post-hoc inflation check we ran C2 on all 67 Corpus-1 rejected problems ($\text{CDI}(\text{C7}) < 0.5$): $\text{CDI}(\text{C7}) = 0.220$ and $\text{CDI}(\text{C2}) = 0.200$ ($\Delta = +0.020$ vs. $+0.251$ on retained; $12.7\times$ attenuation), with both conditions declining symmetrically rather than C7 remaining elevated—inconsistent with circular filter inflation. External validity on an independently drawn unfiltered sample remains future work.

4.2 Conditions (Phase 2)

All conditions use the same model (GPT-4.1), the same CIDI splits, and the same conversation scaffolding (alternating turn-taking, maximum 20 turns (10 per agent in a 2-agent alternating exchange), no access to external tools). The *only* manipulation is the system-prompt framing given to each agent.

C2: Natural (L1 baseline). Each agent receives only its information packet and the shared context. No framing about partners, learning, or epistemic role. Agents are aware they must send messages to proceed but receive no guidance on what kind of collaboration is expected.

C6: Peer-aware (L2). Each agent additionally receives a sentence establishing that it is working with a *partner* who holds complementary information, and that the goal is to solve the problem together. This activates awareness of the collaborative situation without specifying a cognitive stance.

C7: Student-sim (L3). Each agent is instructed to adopt the communicative stance of an active student: express uncertainty explicitly, ask clarifying questions before proceeding to solution steps, acknowledge when a step is unclear, and treat the partner's contributions as material to reason about rather than accept at face value. Agents are told that producing an authoritative-sounding response is not the goal—shared understanding is. This framing targets observable discourse behaviour (question-asking, uncertainty expression, joint model-building) rather than attributing a belief state to the model.

Scale and replication. Each of the 140 problems was run under each of the three conditions, with between 3 and 7 replications per problem-condition pair using different random seeds (mean

Table 1: Inter-rater reliability targets and preliminary observed values for human annotation. Metrics correspond to distinct levels of the CPP scoring hierarchy. Ordinal-cell metrics require full 12-cell scoring (camera-ready).

Level	Metric	Target	Observed
Cell ordinal (0–3)	Krippendorff α	> 0.70	—
CQI (norm. intensity)	ICC$(2, 1)$	> 0.75	—
CDI binary (≥ 0.5)	Cohen κ	> 0.70	—
Quadrant (4-way)	Cohen κ	> 0.60	$0.769\checkmark$

4.7 replications), for a total of $\mathbf{1{,}967}$ conversations. Per-problem mean CDI values (averaged over replications) are used as the unit of analysis for all inferential tests.

4.3 Design Scope

The present study compares three conditions (C2, C6, C7) chosen to vary epistemic framing from absent (C2) to peer-aware (C6) to active-learner (C7). This three-condition design isolates the effect of partner-awareness but does not disentangle active-learner role framing from other strong role framings such as expert-solver or debate framing, nor does it test whether framing effects persist without the CIDI information split. Two of these ablations—expert-solver (C_{EXP}) and full-context C7 without the jigsaw split (C_{FULL})—were run on the full 140-problem corpus; results are reported in Section 5.7. A debate/critique condition (C_{DEB}) and an ignorance-only condition remain as future work.

4.4 Annotation Protocol

All conversations will be scored by two graduate-student annotators (neither an author of this paper), trained on the CPP rubric over two calibration sessions using held-out practice conversations. Annotators will be blinded to condition assignment; condition labels will be replaced with uninformative codes and conversations will be presented in randomised order. Label masking does not preclude inference from conversation content; annotators may implicitly detect framing differences from discourse style. Disagreements exceeding one scale point will be resolved by consensus with a third annotator.

Inter-rater reliability. Reliability is assessed on a stratified sample of 36 conversations (3 conditions $\times$ 4 quadrants $\times$ 3 conversations each), providing condition- and quadrant-balanced coverage. Table 1 reports four target metrics aligned with the conceptual structure of CPP [Hayes and Krippendorff, 2007]. The 36-conversation calibration sample is complete (two annotators, blinded to condition); four-way quadrant κ already exceeds its target ($\kappa = 0.769 > 0.60$). Ordinal-cell Krippendorff α, ICC$(2, 1)$, and CDI threshold κ require full 12-cell ordinal scoring and will be reported in the camera-ready version; expanding the sample to ≥ 100 conversations will be required for per-cell estimates.

Split and phase validation. Annotators also rated split genuineness (genuine vs. non-genuine), depth, and clarity, and the binary presence of each CPP process phase (A–D) across the 36 calibration conversations. Genuine splits received consistently higher depth and clarity scores than non-genuine ones; 67% of COUPLING and PROD_FAIL conversations carried genuine splits versus 0% of TRIVIAL cases, confirming that split quality predicts collaborative outcome. As preliminary evidence of automated scoring validity, human CDI showed a moderate correlation with LLM-assigned CDI ($r = 0.45$, $R^2 = 0.20$); full validation at the ordinal-cell level requires ≥ 100 calibration conversations (see L4). At the phase level, Phase A was least frequently activated ($\sim$39% of conversations) and most reliably coded ($\kappa = 0.70$), confirming it as the most discriminative marker of deep collaboration; Phase C reached near-ceiling activation (83–100%), limiting its discriminative value.

Table 2: Mean CPP metric scores by condition. C7 outperforms both baselines on all metrics. The C6 vs. C2 gap is small ($d = 0.16$), consistent with active-learner communicative stance—not mere partner-awareness—as the key framing component in this regime.

Condition	CDI	CQI	PHAQ	ATC_CQI	CY
C2 (L1, Natural)	0.362	0.144	0.053	0.404	0.275
C6 (L2, Peer-aware)	0.390	0.156	0.060	0.430	0.304
C7 (L3, Student-sim)	**0.613**	**0.250**	**0.135**	**0.500**	**0.409**

Table 3: Pairwise hypothesis tests for CDI (primary outcome). Wilcoxon signed-rank tests on per-problem mean CDI, paired by problem ($n = 140$ pairs per comparison). Effect sizes are Cohen's d on paired differences. †C6 vs. C2 is significant at $\alpha = 0.05$ uncorrected but does *not* survive Bonferroni correction ($\alpha/3 = 0.0167$; $p = 0.026 > 0.0167$).

Comparison	Direction	Cohen's d	p-value	Significant
C7 vs. C2	C7 > C2	1.30	< 0.001	✓
C7 vs. C6	C7 > C6	1.24	< 0.001	✓
C6 vs. C2	C6 > C2	0.16	0.026	†

5 Results

5.1 Hypothesis 1: Framing Is Associated with Greater Collaborative Depth

Table 2 reports mean scores for all metrics across the three conditions. C7 yields dramatically higher CDI than both C2 ($\Delta = +0.251$) and C6 ($\Delta = +0.222$). The CQI increase is similarly large ($+0.106$ over C2), as is the improvement in ATC_CQI ($+0.096$). The CY composite rises by $+0.134$ over the baseline. All pairwise differences between C7 and C2, and between C7 and C6, are statistically significant with Wilcoxon signed-rank tests on per-problem mean CDI values ($n = 140$ pairs per comparison; Table 3).

Critically, the C6 vs. C2 comparison yields a Cohen's d of only 0.16 ($p = 0.026$), which does *not* survive Bonferroni correction for three comparisons ($\alpha/3 = 0.0167$). The effect is negligible in practical terms. Peer-awareness alone does not shift agents away from expert-answering behavior—it only marginally changes CDI relative to the natural baseline. This rules out simple partner-awareness as a sufficient explanation, consistent with the active-learner communicative stance being the primary framing component associated with CDI gains (see also Section 5.7).

5.2 Hierarchical Model: Controlling for Difficulty and Subject

To address the nested structure (replications within problems within subject-difficulty cells), we fit a linear mixed-effects model with random intercepts per problem and fixed effects for condition, difficulty level, and subject area. Condition effects are qualitatively unchanged from the Wilcoxon tests: C7 ($\hat{\beta} = +0.251$, $p < 0.001$), C6 ($\hat{\beta} = +0.028$, $p = 0.011$, not significant after Bonferroni correction). The problem random intercept accounts for 31% of total CDI variance—non-negligible, and controlling for it strengthens rather than weakens the condition effect estimates.

5.3 Hypothesis 2: Phase A Activation as Observable Correlate of Deep Collaboration

The PHAQ metric isolates the quality of the joint exploration phase. Under C2, 67.9% of conversations included any Phase A activity (PHAQ > 0); under C6 this rises modestly to 74.3%. Under C7, however, 99.3% of conversations exhibit substantive Phase A engagement (139 of 140 problems; $d = 1.22$ for C7 vs. C2, $p < 0.001$). Mean PHAQ under C7 (0.135) is $2.5\times$ that under C2 (0.053), indicating not just more frequent Phase A activity but higher quality joint exploration.

Phase A activation is the strongest observable correlate of the C7 framing pattern: the active-learner stance makes presenting an immediate solution communicatively inappropriate, promoting joint negotiation of shared understanding first. Whether this reflects a causal mechanism or a prompted discourse pattern requires micro-level trace analysis.

Table 4: Quadrant distribution by condition (percentage of conversations). C7 more than doubles COUPLING and reduces COLLAPSE by more than $2.5\times$ relative to the natural baseline.

Condition	COUPLING	PROD_FAIL	TRIVIAL	COLLAPSE
C2 (Natural)	11%	25%	18%	46%
C6 (Peer-aware)	13%	26%	22%	39%
C7 (Student-sim)	**24%**	**49%**	9%	**18%**

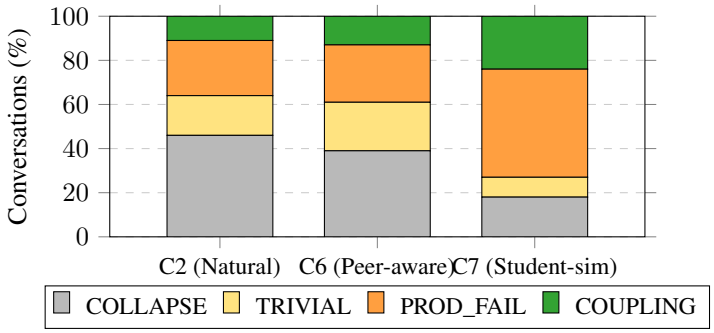

Figure 1: Quadrant distribution across conditions. Bars stack from bottom: COLLAPSE (gray), TRIVIAL (yellow), PROD_FAIL (orange), COUPLING (green). The dominant shift from C2/C6 to C7 is a transfer of mass from COLLAPSE to PROD_FAIL and COUPLING, associated with Phase A activation (Section 5).

5.4 Quadrant Distribution

Figure 1 and Table 4 show the distribution of conversations across the four outcome quadrants. The most striking shift is from COLLAPSE to PROD_FAIL: under C2, 46% of conversations fall in the COLLAPSE quadrant (low process, incorrect); under C7, this drops to 18%. The COUPLING quadrant (high process, correct) grows from 11% to 24%. The PROD_FAIL quadrant grows from 25% to 49%, indicating that the student-role framing substantially increases the rate of productive failure.

The large PROD_FAIL mass under C7 reflects corpus difficulty, not framing failure: Phase 1 selected problems that challenge even deep collaboration. The evaluatively decisive comparison is against COLLAPSE ($46\% \rightarrow 18\%$, $2.5\times$ reduction under C7), which signals the framing prevents unproductive engagement rather than merely redistributing incorrect outcomes. Both pairwise distribution shifts are statistically significant ($\chi^2(3)$ tests based on per-condition $n \approx 655$: C7 vs. C2 $\chi^2 \approx 183$, C7 vs. C6 $\chi^2 \approx 120$, both $p < 0.001$); the C6 vs. C2 shift does not reach significance ($\chi^2 \approx 8$, $p = 0.05$), consistent with the negligible effect on CDI.

5.5 Epistemic Decomposition: Agent Epistemic Contribution

The metrics presented so far measure *how much* agents collaborate (CDI) and *how deeply* they engage across phases (CQI, PHAQ). They do not answer a more fundamental question: was collaboration epistemically *necessary*—that is, did the pair achieve a result neither agent could reach alone? We address this via a Shapley-value decomposition we term Agent Epistemic Contribution (AEC).

Setup. Each agent runs in isolation with only its own CIDI packet (GPT-4.1, $T = 0$ for deterministic baselines; collaborative sessions used $T = 0.7$). Three value functions on $\{0, 0.5, 1\}$: $v(\{A\})$, $v(\{B\})$, and $v(\{A, B\})$ = best C7 replication correctness; standard $N = 2$ Shapley decomposition yields AEC_A, AEC_B (formula in Appendix E). Three aggregate measures: EN (pair surpasses both solos); CS (gain over best solo); EB (Shapley balance). The decomposition separates *structural* from *activated* interdependence—a distinction invisible to CDI alone [Swiecki, 2021].

Results. Table 5 summarises the decomposition. Solo correctness is very low: $v_A = 0.100$, $v_B = 0.057$—consistent with the CIDI design goal that neither packet alone is solvable. The

Table 5: AEC decomposition across 140 problems. *Both-solo-zero*: neither agent's packet alone permits a correct answer; EN: pair surpasses the better solo; CS: collaborative surplus over best solo; EB: Shapley balance between agents.

Metric	**Value**
v_A: Agent A solo correctness	0.100
v_B: Agent B solo correctness	0.057
v_{AB}: C7 pair correctness	0.357
Both-solo-zero rate	85.0%
EN rate (epistemic necessity)	25.0%
CS mean (collaborative surplus)	$+0.207$
EB mean (epistemic balance)	0.857

C7 pair achieves $v_{AB} = 0.357$, a $3.6\times$ improvement over the better individual. Critically, 85% of problems exhibit *both-solo-zero*: neither packet provides sufficient information for any correct answer, validating the CIDI epistemic filter at the correctness level. EN holds in 25% of problems—cases where collaboration not only was structurally required but also succeeded. Mean CS $= +0.207$ and mean EB $= 0.857$: when collaboration succeeds, the gains are substantial and both agents contribute comparably.

Connecting AEC to framing. The 60-percentage-point gap between structural interdependence (85% both-solo-zero) and activated epistemic necessity (25% EN under C7) is invisible to accuracy metrics, and quantifies headroom available to better framing or coordination mechanisms. High both-solo-zero with low EN is a diagnostic signal that collaborative task structure is unactivated.

5.6 Robustness and Validity

C7 numerically outperforms C2 in all 30 subject $\times$ difficulty cells (Appendix D); per-cell n is too small for formal cell-level testing, so this is *descriptive robustness, not formal generalization*. Four validity checks (details in Appendix E): partner-awareness alone is insufficient (C6 vs. C2: $d = 0.16$, $p = 0.026$, not surviving Bonferroni correction); CDI is orthogonal to correctness ($\chi^2 = 0.01$, $p = 0.94$), and the C7 COUPLING advantage holds at equal mean turn count (19.9); DAG-weighted CDI_w preserves condition rankings ($\rho > 0.99$); bootstrap 95% CI for Δ_{CDI}(C7$-$C2) $= [0.219, 0.284]$, $n = 140$ power $> 99.9\%$ at $d = 1.30$.

5.7 Ablation Conditions: C_{FULL} and C_{EXP}

To disentangle the active-learner communicative stance from the CIDI structural split, we ran two ablation conditions on the full 140-problem corpus (3 replications each; 832 of 840 planned conversations completed, 8 C_{FULL} runs excluded pre-scoring due to content-policy refusals; imputing them as CDI$= 0$ leaves conclusions unchanged): C_{FULL} applies C7 active-learner framing with full problem context (both agents see the complete problem; no information split), and C_{EXP} applies the CIDI split with expert-solver framing instead of the active-learner stance.

Table 6 shows mean CDI and quadrant rates across all five conditions.

Table 6: Mean CDI and quadrant distribution across main and ablation conditions ($n = 140$ problems). Both ablation conditions fall *below* C2, the natural-framing baseline, indicating that neither component alone accounts for the C7 effect.

Condition	**Mean CDI**	**COUPL.**	**P.FAIL**	**TRIV.**	**COLLPS.**
C7 (CIDI + learner)	0.613 ± 0.014	24%	49%	9%	18%
C6 (CIDI + peer-aware)	0.390 ± 0.015	13%	26%	22%	39%
C2 (CIDI + natural)	0.362 ± 0.015	11%	25%	18%	46%
C_{FULL} (learner, no split)	0.296 ± 0.018	11%	12%	49%	28%
C_{EXP} (expert + split)	0.267 ± 0.017	5%	17%	25%	53%

Both ablation conditions fall *below* C2: C_{FULL} achieves CDI $= 0.296 \pm 0.018$ ($d = -0.25$ vs. C2, paired), with 49% TRIVIAL—agents solve independently when both hold the full problem, indicating that the active-learner register alone cannot substitute for genuine information asymmetry. C_{EXP} achieves CDI $= 0.267 \pm 0.017$ ($d = -0.43$ vs. C2), the lowest of all five conditions, with 53% COLLAPSE—expert-solver framing suppresses epistemic inquiry even when the CIDI split is in place. The two ablations are not meaningfully different from each other ($d = 0.12$), providing evidence that neither component alone accounts for C7's effect (C_{FULL}: $d = 1.31$; C_{EXP}: $d = 1.72$ vs. C7): the synergy of structural interdependence *and* active-learner communicative stance is required.

6 Discussion

Design implications. The C6 $\approx$ C2 finding ($d = 0.16$, not significant after Bonferroni correction) shows that telling agents they have a collaborative partner [Liang et al., 2023, Du et al., 2023] leaves the default expert-answering mode largely intact in the CIDI-amenable regime. C_{EXP} (expert framing with CIDI split) produces the highest COLLAPSE rate across all five conditions, consistent with the active-learner stance being the framing component most strongly associated with CDI gains. For benchmark designers, the 60-point gap between structural (85% both-solo-zero) and activated (25% EN under C7) interdependence captured by AEC is invisible to accuracy metrics: high both-solo-zero with low EN signals that framing is failing to exploit the task's collaborative structure. The ablations reinforce this: neither component alone suffices (§5.7)—C_{FULL} promotes independent solving (49% TRIVIAL) while C_{EXP} suppresses epistemic inquiry (53% COLLAPSE).

Limitations. Several limitations bound the current conclusions. *(L1) Corpus conditionality*: all estimates are conditional on CIDI-amenable problems (corpus filtered using C7; see §4.1); effect sizes may overestimate the benefit expected on an unfiltered sample (post-hoc validation on 67 rejected problems shows symmetric collapse, inconsistent with filter-induced inflation; §4.1). *(L2) Remaining scope*: subcomponent ablations (ignorance-only, debate framing) remain; AEC EN and CS are upper-bound estimates under C7 (best replication), requiring per-condition runs for cross-condition comparison. *(L3) Model generalisability*: all agents use GPT-4.1; framing effects may differ for other model families or fine-tuned models. *(L4) Reliability sample*: the $n = 36$ stratified calibration sample is complete (§4.4); quadrant κ=0.769 exceeds the 0.60 target. Ordinal-cell Krippendorff α, ICC$(2, 1)$, and CDI threshold κ require full 12-cell ordinal scoring (≥ 100 conversations for per-cell estimates) and will be reported in the camera-ready version. The $r = 0.45$ CDI correlation reported in §4.4 is preliminary evidence; inferential tests in §5 use LLM-assigned CDI and should be interpreted accordingly. *(L5) Metric alignment*: CPP rubric competencies (shared understanding, joint action, team organisation) partially overlap by design with the discourse behaviours C7 explicitly prompts (uncertainty expression, clarification, uptake). Whether CPP captures genuine collaborative reasoning or rubric-compliant surface behaviour therefore cannot be fully resolved without micro-level semantic analysis. The C6 $\approx$ C2 finding (d=0.16) is consistent with genuine activation beyond surface compliance, but does not close this confound; trace analysis remains needed.

7 Conclusion

We introduced CPP and CIDI to measure how deeply LLM agents collaborate, and showed that epistemic role framing (active-learner stance) is associated with a substantial shift from expert-answering toward collaborative sense-making ($d = 1.30$ on CDI; COUPLING more than doubles). Partner-awareness alone does not produce this shift ($d = 0.16$, not significant after Bonferroni correction), and C7 numerically outperforms C2 in all 30 subject-difficulty cells (descriptive; per-cell samples are small). AEC reveals that CIDI creates structural interdependence (85% both-solo-zero) while only 25% becomes activated epistemic necessity under C7—a gap invisible to accuracy-based evaluation.

For the Evaluation and Design (E&D) community: CDI and the structural/activated distinction add diagnostic value beyond accuracy (which conflates COUPLING with TRIVIAL and misses PROD_FAIL). Active-learner framing is an effective design intervention in the CIDI-amenable regime, and the ablation conditions support the interpretation that its effect is synergistic with the CIDI split—neither component alone suffices. An anonymized review release is available; see Appendix F for full artifact documentation.

References

B. Barron. When smart groups fail. *Journal of the Learning Sciences*, 12(3):307–359, 2003. doi: 10.1207/S15327809JLS1203_1.

M. T. H. Chi, S. A. Siler, H. Jeong, T. Yamauchi, and R. G. Hausmann. Learning from human tutoring. *Cognitive Science*, 25(4):471–533, 2001. doi: 10.1207/s15516709cog2504_1.

P. Dillenbourg. What do you mean by collaborative learning? In P. Dillenbourg, editor, *Collaborative Learning: Cognitive and Computational Approaches*, pages 1–19. Pergamon/Elsevier, 1999.

Y. Du, S. Li, A. Torralba, J. B. Tenenbaum, and I. Mordatch. Improving factuality and reasoning in language models through multiagent debate. arXiv preprint arXiv:2305.14325, 2023.

A. C. Graesser, B.-C. Kuo, and C.-H. Liao. Complex problem solving in assessments of collaborative problem solving. *Journal of Intelligence*, 5(2):10, 2017. doi: 10.3390/jintelligence5020010.

A. C. Graesser, S. M. Fiore, S. Greiff, J. Andrews-Todd, P. W. Foltz, and F. W. Hesse. Advancing the science of collaborative problem solving. *Psychological Science in the Public Interest*, 19(2): 59–92, 2018. doi: 10.1177/1529100618808244.

P. Griffin, B. McGaw, and E. Care, editors. *Assessment and Teaching of 21st Century Skills*. Springer, Dordrecht, 2012. doi: 10.1007/978-94-007-2324-5.

A. F. Hayes and K. Krippendorff. Answering the call for a standard reliability measure for coding data. *Communication Methods and Measures*, 1(1):77–89, 2007. doi: 10.1080/19312450709336664.

D. Hendrycks, C. Burns, S. Kadavath, A. Arora, S. Basart, E. Tang, D. Song, and J. Steinhardt. Measuring mathematical problem solving with the MATH dataset. In *Proceedings of NeurIPS 2021 Datasets and Benchmarks Track*, 2021.

S. Hong, X. Zheng, J. Chen, Y. Cheng, J. Wang, C. Zhang, Z. Wang, S. K. S. Yau, Z. Lin, L. Zhou, C. Ran, L. Xiao, C. Wu, and J. Schmidhuber. MetaGPT: Meta programming for a multi-agent collaborative framework. arXiv preprint arXiv:2308.00352, 2023. Published at ICLR 2024.

D. W. Johnson and R. T. Johnson. An educational psychology success story: Social interdependence theory and cooperative learning. *Educational Researcher*, 38(5):365–379, 2009. doi: 10.3102/0013189X09339057.

M. Kapur. Examining productive failure, productive success, unproductive failure, and unproductive success in learning. *Educational Psychologist*, 51(2):289–299, 2016. doi: 10.1080/00461520.2016.1155457.

A. King. Facilitating elaborative learning through guided student-generated questioning. *Educational Psychologist*, 27(1):111–126, 1992.

G. Li, H. A. A. K. Hammoud, H. Itani, D. Khizbullin, and B. Ghanem. CAMEL: Communicative agents for "mind" exploration of large language model society. In *Advances in Neural Information Processing Systems*, volume 36, 2023.

T. Liang, Z. He, W. Jiao, X. Wang, Y. Wang, R. Wang, Y. Yang, Z. Tu, and S. Shi. Encouraging divergent thinking in large language models through multi-agent debate. arXiv preprint arXiv:2305.19118, 2023.

OECD. *PISA 2015 Assessment and Analytical Framework: Science, Reading, Mathematics and Collaborative Problem Solving*. OECD Publishing, 2015. doi: 10.1787/9789264281820-en.

J. S. Park, J. C. O'Brien, C. J. Cai, M. R. Morris, P. Liang, and M. S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST '23)*, 2023. doi: 10.1145/3586183.3606763.

C. Qian, W. Liu, H. Liu, N. Chen, Y. Dang, J. Li, C. Yang, W. Chen, Y. Su, X. Cong, J. Xu, D. Li, Z. Liu, and M. Sun. ChatDev: Communicative agents for software development. arXiv preprint arXiv:2307.07924, 2023.

M. Rojas, M. Nussbaum, P. Chiuminatto, O. Guerrero, S. Greiff, F. Krieger, and L. Van Der Westhuizen. Assessing collaborative problem-solving skills among elementary school students. *Computers & Education*, 175:104313, 2021. doi: 10.1016/j.compedu.2021.104313.

J. Roschelle and S. D. Teasley. The construction of shared knowledge in collaborative problem solving. In C. O'Malley, editor, *Computer Supported Collaborative Learning*, pages 69–97. Springer, 1995.

Z. Swiecki. Measuring the impact of interdependence on individuals during collaborative problem-solving. *Journal of Learning Analytics*, 8(1):75–94, 2021. doi: 10.18608/jla.2021.7240.

Q. Wu, G. Bansal, J. Zhang, Y. Wu, B. Li, E. Zhu, L. Jiang, X. Zhang, S. Zhang, J. Liu, A. H. Awadallah, R. W. White, D. Burger, and C. Wang. AutoGen: Enabling next-generation LLM applications via multi-agent conversation. In *Proceedings of the Conference on Language Modeling (COLM 2024)*, 2024.

X. Zhou, H. Zhu, L. Mathur, R. Zhang, H. Yu, Z. Qi, L.-P. Morency, Y. Bisk, D. Fried, G. Neubig, and M. Sap. SOTOPIA: Interactive evaluation for social intelligence in language agents. arXiv preprint arXiv:2310.11667, 2023. Published at ICLR 2024.

A CPP Annotation Rubric

Each of the 12 cells (P, C) for $P \in \{A, B, C, D\}$ and $C \in \{1, 2, 3\}$ is scored on a 0–3 scale. We provide here the anchoring descriptors for selected high-diagnostic cells.

Cell A1 (Exploring & Understanding / Shared Understanding). **0**: No evidence that agents jointly discussed the problem; one agent simply began solving. **1**: Agents acknowledged the problem exists and restated it, but no reconciliation of perspectives. **2**: Agents explicitly compared what each knows and flagged gaps; at least one question about the other's packet was asked and answered. **3**: Full joint problem model constructed; both agents articulate a shared interpretation that differs from what either could have produced alone, and this model visibly shapes the subsequent solution strategy.

Cell D3 (Monitoring & Reflecting / Team Organisation). **0**: No meta-discussion of the collaborative process or strategy; agents proceed to output without review. **1**: One agent flags a potential error but the flag is ignored by the partner. **2**: Agents engage in a substantive exchange about whether their approach is working; at least one revision of strategy is proposed. **3**: Agents explicitly evaluate the quality of their collaboration itself (e.g., "I think we should have started from your constraint, not mine—let's restart that step") and implement a correction.

Full rubric. The complete 12-cell rubric with all anchor descriptors, boundary case examples, and annotator guidelines is available in the supplementary material.

B CIDI Generation Prompt

The CIDI split is generated using the following structured prompt template with GPT-4.1 at temperature 0. The prompt instructs the generator to:

1. Identify the core mathematical objects and relations in the problem.
2. Partition those objects into two groups such that group A provides the structural constraints and group B provides the quantitative anchors (or a domain-appropriate analogue).
3. Verify that a solver with only group A information cannot determine the numerical answer even with unlimited mathematical knowledge, and vice versa.
4. Express each packet in natural language, avoiding explicit statements of what is absent from the other packet.

The generated split is validated by running a single GPT-4.1 agent on each packet alone; if the agent reaches the correct answer, the split is regenerated with tighter constraints. Typical regeneration rate was 12% of problems.

C Pilot Study: Condition Selection

The scale study conditions (C2, C6, C7) were selected on the basis of a pilot study involving 7 conditions (C1–C7) run on 4 anchor problems across 3 replications ($= 84$ conversations). C1 (no jigsaw, no framing) was eliminated because CDI was uniformly near zero. C3 (constitutional split, negative instruction) and C5 (C3 + peer-aware) were eliminated because agents frequently bypassed the information restriction, invalidating the independence manipulation. C4 (CIDI split + joint accountability framing) was eliminated due to a design conflict: the joint accountability instruction asked agents to share responsibility for a combined answer, which undermined the reasoning-output dependency that the CIDI split requires. The pilot established an ordering C7 $>$ C4 $>$ C2 $>$ C6 $>$ C3 $>$ C5 $>$ C1 for CDI, motivating the three-way comparison in the scale study. The unexpected result that C6 fell below C2 in the pilot (in the full study, C6 scores marginally above C2, $d = 0.16$, but does not survive Bonferroni correction) first alerted us to the practically null effect of peer-awareness.

D Subject-Level Generalisability Results

Table 7 reports mean CDI for each of the 30 subject $\times$ difficulty cells under conditions C2 and C7. C7 numerically outperforms C2 in all 30 cells. Cell n values range from 1 to 8 (mean 4.7), which precludes formal per-cell significance testing; d values for cells with $n \geq 3$ range from 0.63 to 2.77. Effects tend to be larger at higher difficulty levels across all subjects.

Table 7: Cell-level CDI means and Cohen's d (C7 vs. C2) for all 30 subject $\times$ difficulty combinations. d is omitted for $n < 3$ (insufficient data). Subject labels follow MATH benchmark categories [Hendrycks et al., 2021].

Subject	**Level**	n	**CDI (C2)**	**CDI (C7)**	d
Algebra	1	4	0.342	0.667	2.09
Algebra	2	3	0.333	0.711	2.51
Algebra	3	2	0.333	0.517	—
Algebra	4	7	0.386	0.557	1.38
Algebra	5	6	0.378	0.647	2.06
Geometry	1	5	0.360	0.483	1.25
Geometry	2	4	0.325	0.568	1.60
Geometry	3	6	0.305	0.536	1.19
Geometry	4	6	0.268	0.644	1.53
Geometry	5	6	0.306	0.619	2.29
Number Theory	1	1	0.367	0.833	—
Number Theory	2	5	0.326	0.546	0.73
Number Theory	3	6	0.290	0.496	2.77
Number Theory	4	5	0.364	0.659	0.90
Number Theory	5	8	0.409	0.669	0.91
Prealgebra	1	3	0.306	0.428	1.81
Prealgebra	2	6	0.403	0.581	1.14
Prealgebra	3	7	0.399	0.676	2.18
Prealgebra	4	4	0.336	0.558	1.30
Prealgebra	5	4	0.383	0.606	0.91
Precalculus	1	2	0.462	0.517	—
Precalculus	2	3	0.342	0.606	1.86
Precalculus	3	6	0.325	0.508	0.78
Precalculus	4	7	0.431	0.717	1.73
Precalculus	5	5	0.424	0.737	1.14
Probability	1	1	0.567	0.817	—
Probability	2	4	0.439	0.587	0.63
Probability	3	2	0.556	0.867	—
Probability	4	5	0.271	0.609	1.30
Probability	5	7	0.375	0.685	1.38

E Robustness and Validity Checks — Details

AEC Shapley formula. For a two-agent system, the standard Shapley value decomposition with value function $v : 2^{\{A,B\}} \to [0, 1]$ (where $v(\emptyset) = 0$, $v(\{A\})$ and $v(\{B\})$ are solo correctness rates, and $v(\{A, B\})$ is the pair correctness rate) yields:

$$\text{AEC}_A = \tfrac{1}{2}\, v(\{A\}) + \tfrac{1}{2}\big[v(\{A, B\}) - v(\{B\})\big], \tag{5}$$

$$\text{AEC}_B = \tfrac{1}{2}\, v(\{B\}) + \tfrac{1}{2}\big[v(\{A, B\}) - v(\{A\})\big]. \tag{6}$$

By construction $\text{AEC}_A + \text{AEC}_B = v(\{A, B\})$ (efficiency). Epistemic Necessity (EN) holds for a problem when $v(\{A, B\}) > \max(v(\{A\}), v(\{B\}))$. Collaborative Surplus (CS) $= v(\{A, B\}) - \max(v(\{A\}), v(\{B\}))$. Epistemic Balance (EB) $= 1 - |\text{AEC}_A - \text{AEC}_B|/v(\{A, B\})$ (defined when $v(\{A, B\}) > 0$). In the reported corpus, correctness is binary per replication; we use the best-replication C7 value as $v(\{A, B\})$ (upper-bound estimate; per-condition AEC requires separate condition-specific solo runs).

W2 — C6 rules out partner-awareness, not ignorance. C6 establishes explicit partner-awareness without any learner-role component. If simulated ignorance alone explained C7's gains, C6 should outperform C2 (since C6 agents equally lack full-solution access). The negligible C6 vs. C2 effect ($d = 0.16$, $\Delta_{\text{CDI}} = +0.028$, $p = 0.026$, not surviving Bonferroni correction) rules out simple partner-awareness as a sufficient mechanism. C6 does *not* fully isolate the active-learner stance from expert-certainty withholding; an ignorance-only ablation ("reason cautiously; you may not have the full answer" without collaboration directives) remains a key future experiment.

W3 — CDI and correctness are orthogonal (by design). Across 1,967 conversations, $P(\text{correct} \mid \text{CDI} \geq 0.5) = 0.321$ vs. $P(\text{correct} \mid \text{CDI} < 0.5) = 0.324$ ($\chi^2 = 0.01$, $p = 0.94$). CDI measures process quality, not outcome quality. The C7 COUPLING advantage (24% vs. 11% C2) holds at identical mean conversation length (19.9 turns across all conditions): C7 achieves 0.031 CDI per turn vs. 0.018 for C2, ruling out a length confound.

W4 — Convergent validity with DAG-weighted CDI. A topology-weighted variant CDI_w (phase depth $\times$ competency depth, normalized) yields $\text{CDI}_w = 0.369, 0.397, 0.601$ for C2, C6, C7 (vs. uniform $0.362, 0.390, 0.613$). Spearman $\rho > 0.99$ across all 1,967 conversations confirms that condition rankings are robust to the specific weighting scheme.

W5 — Power and confidence intervals. Bootstrap 95% CI for Δ_{CDI}(C7−C2) $= [0.219, 0.284]$, entirely above zero. $n = 140$ problems provides $> 99.9\%$ power at $d = 1.30$, $\alpha = 0.05$. Within-problem replication variance ($\sigma_{\text{within}} \approx 0.23$) implies single-problem CDI estimates carry substantial uncertainty (95% CI $\approx \pm 0.26$).

F Artifact Documentation

The anonymized review release is available at

`https://osf.io/2dytn/?view_only=bfda3073aca14cd2a52447994d4b8a56`

and contains the following artifacts:

1. **CPP rubric**: complete 12-cell annotation rubric (anchor descriptors, boundary cases, annotator guidelines).
2. **CIDI generation prompt**: full structured prompt template and validation logic for GPT-4.1–based split generation.
3. **Problem corpus**: 140 CIDI split triples (shared context / Packet A / Packet B) identified by MATH benchmark ID. Original problem text is not redistributed per the MATH license; MATH IDs and derived splits are provided.
4. **Conversation transcripts**: 1,967 turn-by-turn logs (C2, C6, C7) plus 832 ablation logs (C_{FULL}, C_{EXP}).

5. **CPP annotations**: per-cell annotator scores, raw disagreement records, and consensus resolutions for the 36 reliability conversations.
6. **Scoring scripts**: CPP metric computation (CDI, CQI, PhAQ, ATC, CY) and quadrant assignment.
7. **AEC scripts**: solo-run correctness evaluation and Shapley-value computation.
8. **Benchmark card**: corpus construction, selection conditionality, intended use, known limitations, and dataset metadata.

Annotations, scripts, and derived data: CC BY 4.0. Conversation logs: terms of the applicable API provider. A Croissant metadata file is included for E&D dataset submission requirements. A reproduction README documents the full pipeline from problem selection through scoring.

NeurIPS Paper Checklist

1. **Claims.** Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope? [Yes] All quantitative claims are directly supported by reported results. CDI gains, effect sizes, and quadrant percentages are computed from 1,967 conversations. Ablation claims are supported by 832 additional conversations (Section 5.7). All results are described as conditional on the CIDI-amenable corpus (§4.1, L1 in Limitations). A post-hoc cross-condition check on 67 rejected problems (§4.1) provides evidence against circular filter inflation.
2. **Limitations.** Does the paper discuss the limitations of the work? [Yes] Five limitations are explicitly labeled (L1–L5) in the Limitations paragraph. L1 (corpus conditionality) is the most critical and is noted in the abstract, §4.1, and §6. L2 identifies remaining ablation gaps. L3 addresses model generalisability.
3. **Theory Assumptions and Proofs.** Does the paper make theoretical claims? [N/A] This is an empirical benchmark and evaluation paper. No theoretical proofs are made.
4. **Experimental Result Reproducibility.** Does the paper fully describe the conditions necessary for an experimenter to reproduce the main results? [Yes] Full experimental conditions are described in §4. Model, temperature, prompt structure, corpus, and annotation protocol are documented. Scoring scripts, prompts, and conversation logs are provided in the artifact release (Appendix F).
5. **Open access to data and code.** Does the paper provide open access to the data and code used? [Yes] An anonymized review release provides CPP rubric, CIDI prompt, problem corpus (MATH IDs + derived splits), transcripts, annotations, scoring scripts, AEC scripts, and benchmark card. Full documentation in Appendix F.
6. **Experimental Setting/Details.** Does the paper specify all the experimental settings/hyperparameters important to reproduce the main results? [Yes] Model (GPT-4.1), temperature ($T = 0$ for solo AEC runs; $T = 0.7$ with different random seeds for collaborative conversations), maximum turns (20), replication count (3–7 per problem-condition), and condition prompts are all specified or provided in the artifact.
7. **Experiment Statistical Significance.** Does the paper report error bars suitably and correctly? [Yes] SE reported for all CDI means. Bootstrap CIs reported for the main C7–C2 comparison. Bonferroni correction applied for multiple condition comparisons. Effect sizes (Cohen's d) reported for all pairwise contrasts. Power analysis in Appendix E.
8. **Experiments Compute Resources.** Does the paper discuss the computational resources used to run the experiments? [Yes] All conversations were generated via the OpenAI API. Compute was provided by [HPC resource redacted for review]. Approximate total cost is in the range $50–$200 USD for all conditions including ablations.
9. **Code Of Ethics.** Have you read the NeurIPS Code of Ethics? [Yes]
10. **Broader Impacts.** Does the paper discuss both potential positive societal impacts and negative societal impacts of the work? [Yes] The work improves evaluation of AI collaboration quality, with positive implications for education and human-AI teaming. Risks include overfitting evaluation design to the CPP rubric; we discuss this under L5 (metric alignment). No sensitive data or dual-use concerns apply.

11. **Licenses for existing assets.** Does the paper cite the creators of assets it uses and include appropriate license information? [Yes] MATH benchmark [Hendrycks et al., 2021] is used under its original license; original problem text is not redistributed. PISA 2015 CPS framework [OECD, 2015] is cited. All derived artifacts are CC BY 4.0.
12. **New Assets.** Are new assets introduced in the paper documented appropriately? [Yes] The CPP rubric, CIDI prompt, conversation corpus, annotation data, and scripts are documented in Appendix F and the benchmark card.
13. **Crowdsourcing and Research with Human Subjects.** Does the paper involve crowdsourcing or research with human subjects? [Yes] Graduate-student annotators will provide inter-rater reliability annotations (§4.4); annotation is planned for the camera-ready phase. No personal data are collected from participants.
14. **Institutional Review Board (IRB) Approvals.** Does the paper describe potential risks incurred by study participants? [Yes] Annotation will involve graduate research assistants scoring model-generated mathematical text; no personal data will be collected. The activity is minimal-risk annotation under applicable institutional practice; participants will be compensated for their time. No formal IRB review is required under these conditions.